

Intro to Deep Learning

Big Data y Machine Learning para Economía Aplicada

Ignacio Sarmiento-Barbieri

Universidad de los Andes


Deep Learning: Intro

- ▶ Linear Models may miss the nonlinearities that best approximate $f^*(x)$
- ▶ Neural networks are simple models.
- ▶ The model has **linear combinations** of inputs that are passed through **nonlinear activation functions** called nodes (or, in reference to the human brain, neurons).

Agenda

- 1 Recap: SNN
 - Activation Functions
 - Output Functions
- 2 Training the network
- 3 Architecture Design
 - Deep Neural Networks

Single Layer Neural Networks



Single Layer Neural Networks

- ▶ NN are made of **linear combinations** of inputs that are passed through **nonlinear activation functions**
- ▶ The NN model has the form

$$f(X) = f \left[\beta_0 + \sum_{k=1}^K \beta_k A_k \right] \quad (1)$$

$$= f \left[\beta_0 + \sum_{k=1}^K \beta_k g \left(w_{k0} + \sum_{j=1}^p w_{kj} X_j \right) \right] \quad (2)$$

- ▶ where
 - ▶ $g(.)$ is a activation function, the nonlinearity of $g(.)$ is **key**
 - ▶ f is the output layer of the network
- ▶ both are prespecified

Agenda

- 1 Recap: SNN
 - Activation Functions
 - Output Functions
- 2 Training the network
- 3 Architecture Design
 - Deep Neural Networks

Neural Networks: Activation Functions

- ▶ $\text{Sigmoid}(x) = \frac{1}{1+\exp(-x)}$
- ▶ $\text{ReLU}(x) = \max\{x, 0\}$
- ▶ Among others ([see more here](#))
- ▶ Hidden unit design remains an active area of research, and many useful hidden unit types remain to be discovered

Agenda

- ① Recap: SNN
 - Activation Functions
 - Output Functions
- ② Training the network
- ③ Architecture Design
 - Deep Neural Networks

Output Functions

- ▶ The choice of output unit is related to the problem at hand
 - ▶ Regression
 - ▶ Classification
 - ▶ Binary
 - ▶ Multiclass

Agenda

- ① Recap: SNN
 - Activation Functions
 - Output Functions
- ② Training the network
- ③ Architecture Design
 - Deep Neural Networks

Training the network

- El objetivo es


$$\hat{f} = \operatorname{argmin}_f \left\{ \sum_{i=1}^n L(y, f(X; \Theta)) \right\} \quad (3)$$

- SNN

$$f(X, \beta, w) = f \left[\beta_0 + \sum_{k=1}^K \beta_k g \left(w_{k0} + \sum_{j=1}^p w_{kj} X_j \right) \right] \quad (4)$$

Training the network

Example: House Prices



► Equations

► Hidden Layer sigmoid (logistic):

$$\text{► } A_1 = \sigma(w_{11} \cdot X_1 + w_{12} \cdot X_2 + w_{10})$$

$$\text{► } A_2 = \sigma(w_{21} \cdot X_1 + w_{22} \cdot X_2 + w_{20})$$

► Output Layer, identity output function:

$$\text{► } \hat{y}_i = \beta_0 + \beta_1 \cdot A_1 + \beta_2 \cdot A_2$$


► Loss Function \Rightarrow MSE: $\frac{1}{n} \sum_i^n (y_i - \hat{y}_i)^2$

Training the network


- El objetivo es

$$\hat{f} = \operatorname{argmin}_{w, \beta} \left\{ \sum_{i=1}^n L(y, f \left[\beta_0 + \sum_{k=1}^K \beta_k g \left(w_{k0} + \sum_{j=1}^p w_{kj} X_j \right) \right]) \right\} \quad (5)$$

Training the network




Training the network



Training the network

Example: House Prices



► Equations

► Hidden Layer sigmoid (logistic):

$$\text{► } A_1 = \sigma(w_{11} \cdot X_1 + w_{12} \cdot X_2 + w_{10})$$

$$\text{► } A_2 = \sigma(w_{21} \cdot X_1 + w_{22} \cdot X_2 + w_{20})$$

► Output Layer, identity output function:

$$\text{► } \hat{y}_i = \beta_0 + \beta_1 \cdot A_1 + \beta_2 \cdot A_2$$

► Loss Function \Rightarrow MSE: $\frac{1}{n} \sum_i^n (y_i - \hat{y}_i)^2$

How does backpropagation work?

Updating a single weight

- For simplicity let's focus on updating w_{11}

Training the network

Updating a weight in the output layer

- Now let's update one of the weights from the hidden layer to the output layer, β_1

Training the network


Batch Gradient Descent

- ▶ Notice that this formula involves calculations over the full data set, at each Gradient Descent step!
- ▶ This is why the algorithm is called Batch Gradient Descent: it uses the whole batch of training data at every step.
- ▶ As a result it is terribly slow on very large data sets.

Training the network

Stochastic Gradient-Based Optimization

- Stochastic Gradient Descent just picks a random observation at every step and computes the gradients based only on that single observation.




Training the network

Mini-batch Gradient Descent

- ▶ Batch Gradient Descent involves calculations over the full data set
- ▶ Stochastic Gradient Descent just picks one random observation at every step
- ▶ At each step, mini- batch GD computes the gradients on small random sets of observations called mini- batches.

Training the network



Training the network

SGD + Momentum, Polyak, 1964

$$v' = \gamma v - \epsilon \cdot \nabla_{\theta} \mathcal{L}(\theta)$$

$$\theta' = \theta + v'$$

- ▶ Agrega una "inercia" al gradiente.
- ▶ Permite acumular dirección → suaviza la trayectoria.
- ▶ Ayuda a escapar de mínimos poco profundos.

Training the network

RMSProp, Tieleman and Hinton, 2012

$$E[\nabla_{\theta}\mathcal{L}(\theta)^2] = \rho E[\nabla_{\theta}\mathcal{L}(\theta)^2]_{iter-1} + (1 - \rho) \nabla_{\theta}\mathcal{L}(\theta)^2$$
$$\theta' = \theta - \frac{\epsilon}{\sqrt{E[\nabla_{\theta}\mathcal{L}(\theta)^2] + \eta}} \cdot \nabla_{\theta}\mathcal{L}(\theta)_t$$

Training the network

Adam, Kingma and Ba, 2014

$$m = \beta_1 m_{iter-1} + (1 - \beta_1) \nabla_{\theta} \mathcal{L}(\theta)$$

$$v = \beta_2 v_{iter-1} + (1 - \beta_2) \nabla_{\theta} \mathcal{L}(\theta)^2$$

$$\hat{m} = \frac{m}{1 - \beta_1^{iter}}, \quad \hat{v} = \frac{v}{1 - \beta_2^{iter}}$$

$$\theta' = \theta - \epsilon \cdot \frac{\hat{m}}{\sqrt{\hat{v} + \eta}}$$

- Combina Momentum + RMSProp.
- Corrige sesgo en los primeros pasos.
- Muy popular por su robustez y rapidez.

Training the network


Which one to use? Zhou et al., 2020

- ▶ Aunque Adam converge más rápido, SGD suele encontrar soluciones que generalizan mejor.
- ▶ Zhou et al., 2020 evalúa ambos métodos
- ▶ SGD: Puede
 - ▶ Escapar más fácilmente de mínimos locales.
 - ▶ Alcanzar soluciones que generalizan mejor.
- ▶ Adam: quedar atrapado en mínimos subóptimos.

Conclusión: El ruido inherente al SGD actúa como una forma de regularización que favorece la generalización.

Training the network

Example: MNIST



Agenda

- 1 Recap: SNN
 - Activation Functions
 - Output Functions
- 2 Training the network
- 3 Architecture Design
 - Deep Neural Networks

Architecture Design

- ▶ A key design consideration for neural networks is determining the architecture.
- ▶ The word architecture refers to the overall structure of the network: how many units it should have and how these units should be connected to each other.
- ▶ The universal approximation theorem (Hornik et al., 1989; Cybenko, 1989) guarantees that regardless of what function we are trying to learn, a sufficiently large MLP will be able to represent this function.

Architecture Design

- ▶ A key design consideration for neural networks is determining the architecture.
- ▶ The word architecture refers to the overall structure of the network: how many units it should have and how these units should be connected to each other.
- ▶ The universal approximation theorem (Hornik et al., 1989; Cybenko, 1989) guarantees that regardless of what function we are trying to learn, a sufficiently large MLP will be able to represent this function.
- ▶ However, learning can fail for two different reasons.
 - 1 The optimization algorithm used for training may not be able to find the value of the parameters that corresponds to the desired function.
 - 2 The training algorithm might choose the wrong function as a result of overfitting

Architecture Design

- ▶ Using deeper models can reduce the number of units required to represent the desired function and can reduce the amount of generalization error.
- ▶ The ideal network architecture for a task must be found via experimentation guided by monitoring the validation set error

Agenda

- ① Recap: SNN
 - Activation Functions
 - Output Functions
- ② Training the network
- ③ Architecture Design
 - Deep Neural Networks


Architecture Design

Multilayer Neural Networks

- ▶ Modern neural networks typically have more than one hidden layer, and often many units per layer.
- ▶ In theory a single hidden layer with a large number of units has the ability to approximate most functions.
- ▶ However, the learning task of discovering a good solution is made much easier with multiple layers each of modest size.

Architecture Design

Multilayer Neural Networks



Architecture Design

Network Tuning

- ▶ Training networks requires a number of choices that all have an effect on the performance:
 - ▶ The number of hidden layers,
 - ▶ The number of units per layer
 - ▶ Details of stochastic gradient descent.
 - ▶ Regularization of parameters
- ▶ This is an active research area that involves a lot of trial and error, and overfitting is a latent danger at each step.

Training the network

Example: MNIST

